

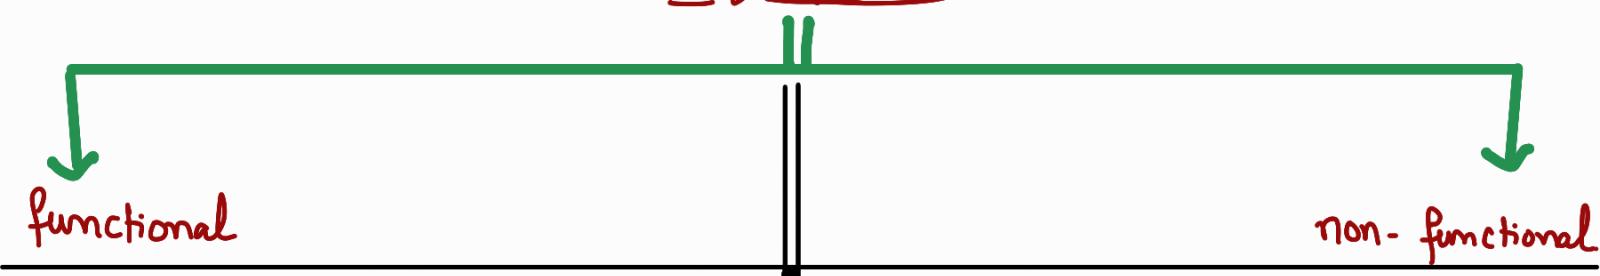
How to
approach

* Generally open ended problems are given.

* So, we will dig into the problem, like

- SCOPING**
- {
(i) whether we wanted to high level design of the full web-site
(ii) Or, how a particular functionality like (liking a tweet) will work.

Requirements



- * How we want the application to function? What should happen when someone likes a tweet?
How to generate the user's feed?
- * Discussed so that we will design up only what's required, neither what I wanted
- * Used to check how careful we are? or we just make assumptions without clarifying the statement

- * Mainly about Scalability
 - ↳ like we have 100M daily users. So, how we should handle these many requests?
- * Throughput
- * Storage Capacity of our D.B.
- * Performance
 - * latency
 - * D.B partitions.
 - ↳ like read < 1 sec
 - write < 2 second
- * Availability.

* we are not required to over-engineer our solution. Because often a particular threshold (5 g's), the improvement is of 5 mins, but cost is going to be really high.

* Tradeoff :- To let non-important things go, then losing the Super-important functionality..

Back of the envelope calculations

* In the interview, we do not have time for the calculation of exact no. of

* Let us understand it with an example :-

Suppose we have a user base of 100M out of which every user on an average reads about 100 tweets daily and 10% of these actually write tweets.

$$\Rightarrow \text{Total reads} \Rightarrow 100 \text{ M} \times 100 \text{ reads} = 10 \text{ billion.}$$

$$\text{Total writes} \Rightarrow 10\% \text{ of } 100 \text{ M} = 10 \text{ million}$$

$$\text{Throughput}_{\text{read}} = \frac{10b}{\text{Total second in a day}} = \frac{10 \times 10^9}{60 \times 60 \times 24}$$

$$\approx \frac{10 \text{ billion}}{100 \text{ k.}} = 10^5 \text{ reads / second.}$$

$$\text{Throughput}_{\text{write}} = 10^2 \text{ write / seconds.}$$

* for simplicity, I will ignore the requests of authentication & payments.

2020

■
1ns

■
L1 cache reference: 1ns

■
Main memory reference: 100ns
1,000ns ≈ 1μs

■
Branch mispredict: 3ns
L2 cache reference: 4ns

■
Mutex lock/unlock: 17ns
10,000ns ≈ 10μs = ■

■
Compress 1KB wth Zippy: 2,000ns
≈ 2μs

■
Read 1,000,000 bytes sequentially
from memory: 3,000ns ≈ 3μs

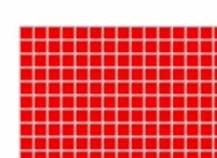
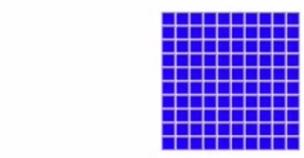
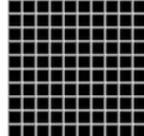
■
Round trip in same datacenter:
500,000ns ≈ 500μs

■
Send 2,000 bytes over commodity
network: 44ns

■
SSD random read: 16,000ns ≈ 16μs

■
Disk se

■
100ns = ■



■
Read 1,
from Si
Packet
Netherl
150ms

